

NOTES 12: μ OCAML COMPILER

Cameron Moy

Tuesday July 10th

1 λ syntax and semantics

Here is a grammar for the λ calculus.

$$T \rightarrow T\ T \mid \lambda X.T \mid X$$

$$X \rightarrow a \mid b \mid \dots \mid z \mid aa \mid \dots$$

And here are the semantics.

$$(\alpha) \frac{}{(\lambda x.t) \Rightarrow_r (\lambda y.t[x \mapsto y])}$$

$$(\beta) \frac{}{(\lambda x.t_1)\ t_2 \Rightarrow_r t_1[x \mapsto t_2]}$$

$$(3) \frac{t_1 \Rightarrow_r t'_1}{t_1\ t_2 \Rightarrow_r t'_1\ t_2} \quad (4) \frac{t_2 \Rightarrow_r t'_2}{t_1\ t_2 \Rightarrow_r t_1\ t'_2} \quad (5) \frac{t \Rightarrow_r t'}{(\lambda x.t) \Rightarrow_r (\lambda x.t')}$$

The semantics define a reduction relation, not a function, $\Rightarrow_r: T \rightarrow T$. This style of semantics is called small-step. The fact \Rightarrow_r is a relation means our semantics are now non-deterministic.

Rule (α) is α -conversion. Rule (β) is β -reduction. Rules (3) through (5) don't actually do anything, but allow us to apply the others in any sub-expression. Rules that do the work, like (α) and (β) , are called computation rules. Rules

that allow computation to happen in sub-expressions, like (3) through (5), are called congruence rules.

Example. Reduce $(\lambda x.x (\lambda y.x y)) y$ to normal form.

Proof. We could do this proof as a series of trees, but this is a lot of writing. Typically we just show the series of \Rightarrow_r directly, underlining the redex.

$$\begin{aligned} (\lambda x.(\lambda y.x y) x) y &\Rightarrow_r^\alpha (\lambda x.(\lambda z.x z) x) y \\ &\Rightarrow_r^\beta (\lambda z.y z) y \\ &\Rightarrow_r^\beta y y \end{aligned}$$

Notice that we are required to α convert in this case, otherwise we are prone to variable capture. \square

Since \Rightarrow_r is a relation (and hence non-deterministic), there is actually another way to do this proof.

Proof.

$$\begin{aligned} (\lambda x.(\lambda y.x y) x) y &\Rightarrow_r^\beta (\lambda x.x x) y \\ &\Rightarrow_r^\beta y y \end{aligned}$$

\square

Notice that we get the exact same answer. This is actually a guarantee.

Theorem (Church-Rosser). If $t \Rightarrow_r^* u$ and $t \Rightarrow_r^* v$ (where \Rightarrow_r^* is an arbitrary series of \Rightarrow_r steps) and u and v are normal forms, then $u = v$.

This property is known as confluence and tells us that terms in our calculus have a unique normal form, no matter what reduction strategy you take. Notice the condition here, confluence only holds if we end up at a normal form. Not all reduction strategies get us to the normal form!

Example. Let $\Omega \triangleq (\lambda x.x\ x) (\lambda x.x\ x)$. What is $((\lambda x.y)\ \Omega)$ under CBV and CBN?

2 Reduction strategies

This is all well-and-good, but non-determinism is bad. If we want to write a λ calculus interpreter we need to define a deterministic strategy. We can describe such a strategy informally.

1. The call-by-value strategy requires that an argument must be evaluated before a β reduction can happen.
2. The call-by-name strategy states that β reductions happen without evaluating the argument at all.

However, we can use our other style of semantics, big-step semantics, to formally describe a deterministic reduction strategy. Here is such a CBV semantics.

$$\begin{array}{c}
 (1) \frac{}{\lambda x.t \Downarrow_v \lambda x.t} \\
 (2) \frac{t_1 \Downarrow_v \lambda x.t_{12} \quad t_2 \Downarrow_v v_2 \quad t_{12}[x \mapsto v_2] \Downarrow_v v}{t_1\ t_2 \Downarrow_v v}
 \end{array}$$

3 It was λ the whole time

Pause for a moment. Doesn't this look familiar? NanOCaml is really just the λ calculus in disguise. A couple differences are worth noting.

1. Our syntax had **fun** instead of λ , an \rightarrow instead of $.$, and extra parentheses (to make parsing easier). In the actual λ calculus we omit most parentheses and use two conventions to disambiguate: λ bodies extend as far as possible, and application left associates.

2. The semantics we used were not the non-deterministic small-step kind, but were the deterministic big-step kind. We choose our semantics to implement a CBV reduction strategy. We could've made a different choice and would've gotten a different semantics.