

NOTES 12: μ OCAML COMPILER

Cameron Moy

Tuesday July 10th

1 λ syntax and semantics

Here is a grammar for the λ calculus.

$$T \rightarrow T\ T \mid \lambda X.T \mid X$$

$$X \rightarrow a \mid b \mid \dots \mid z \mid aa \mid \dots$$

And here are the semantics.

$$(\alpha) \frac{}{(\lambda x.t) \Rightarrow_r (\lambda y.t[x \mapsto y])}$$

$$(\beta) \frac{}{(\lambda x.t_1)\ t_2 \Rightarrow_r t_1[x \mapsto t_2]}$$

$$(3) \frac{t_1 \Rightarrow_r t'_1}{t_1\ t_2 \Rightarrow_r t'_1\ t_2} \quad (4) \frac{t_2 \Rightarrow_r t'_2}{t_1\ t_2 \Rightarrow_r t_1\ t'_2} \quad (5) \frac{t \Rightarrow_r t'}{(\lambda x.t) \Rightarrow_r (\lambda x.t')}$$

The semantics define a reduction relation, not a function, $\Rightarrow_r: T \rightarrow T$. This style of semantics is called small-step. The fact \Rightarrow_r is a relation means our semantics are now non-deterministic.

Rule (α) is α -conversion. Rule (β) is β -reduction. Rules (3) through (5) don't actually do anything, but allow us to apply the others in any sub-expression. Rules that do the work, like (α) and (β) , are called computation rules. Rules

that allow computation to happen in sub-expressions, like (3) through (5), are called congruence rules.

Example. Reduce $(\lambda x.x (\lambda y.x y)) y$ to normal form.

Proof. We could do this proof as a series of trees, but this is a lot of writing. Typically we just show the series of \Rightarrow_r directly, underlining the redex.

$$\begin{aligned} (\lambda x.(\lambda y.x y) x) y &\Rightarrow_r^\alpha (\lambda x.(\lambda z.x z) x) y \\ &\Rightarrow_r^\beta (\lambda z.y z) y \\ &\Rightarrow_r^\beta y y \end{aligned}$$

□

Since \Rightarrow_r is a relation (and hence non-deterministic), there is actually another way to do this proof.

Proof.

$$\begin{aligned} (\lambda x.(\lambda y.x y) x) y &\Rightarrow_r^\beta (\lambda x.x x) y \\ &\Rightarrow_r^\beta y y \end{aligned}$$

□

Notice that we get the exact same answer. This is actually a guarantee.

Theorem (Church-Rosser). If $t \Rightarrow_r^* u$ and $t \Rightarrow_r^* v$ (where \Rightarrow_r^* is an arbitrary series of \Rightarrow_r steps) and u and v are normal forms, then $u = v$.

This property is known as confluence and tells us that terms in our calculus have a unique normal form, no matter what reduction strategy you take.

2 Reduction strategies

This is all well-and-good, but non-determinism is bad. If we want to write a λ calculus interpreter we need to define a deterministic strategy. We can describe such a strategy informally.

1. The call-by-value strategy requires that an argument must be evaluated before a β reduction can happen.
2. The call-by-name strategy states that β reductions happen without evaluating the argument at all.

However, we can use our other style of semantics, big-step semantics, to formally describe a deterministic reduction strategy. Here is such a CBV semantics.

$$\begin{array}{c} (1) \frac{}{(\mathbf{fun} \ x \rightarrow t) \Downarrow_v (\mathbf{fun} \ x \rightarrow t)} \\ (2) \frac{t_1 \Downarrow_v (\mathbf{fun} \ x \rightarrow t_{12}) \quad t_2 \Downarrow_v v_2 \quad t_{12}[x \mapsto v_2] \Downarrow_v v}{(t_1 \ t_2) \Downarrow_v v} \end{array}$$

3 It was λ the whole time

Pause for a moment. Doesn't this look familiar?

You should have realized by now that last week we implemented a CBV λ calculus interpreter. We just called it “NanOCaml”. Our syntax had **fun** instead of λ , an \rightarrow instead of $.$, and extra parentheses (to make parsing easier), but really that's it!